

A PROJECT REPORT

on

“CROP YIELD PREDICTION”

Submitted to

KIIT Deemed to be University

In Partial Fulfilment of the Requirement for the Award of

BACHELOR’S DEGREE IN

INFORMATION TECHNOLOGY

BY

ABHIJEET KUMAR	2106002
AMAN KUMAR	2106011
SHOUVIK GHOSH	2106067
TUSHAR BHATACHARYA	2106079
SUKHARANAJAN JANA	2106269

**UNDER THE GUIDANCE OF
MS. IPSITA PAUL**



SCHOOL OF COMPUTER ENGINEERING

KALINGA INSTITUTE OF INDUSTRIAL TECHNOLOGY

BHUBANESWAR, ODISHA - 751024

April 2024

PROJECT REPORT

on

“CROP YIELD PREDICTION”

Submitted to

KIIT Deemed to be University

In Partial Fulfilment of the Requirement for the Award of

**BACHELOR’S DEGREE IN
INFORMATION TECHNOLOGY**

BY

UNDER THE GUIDANCE OF

MS. IPSITA PAUL



SCHOOL OF COMPUTER ENGINEERING

KALINGA INSTITUTE OF INDUSTRIAL TECHNOLOGY

BHUBANESWAE, ODISHA -751024

April 2024

KIIT Deemed to be University

School of Computer Engineering

Bhubaneswar, ODISHA 751024



CERTIFICATE

This is certified that the project entitled

“CROP YIELD PREDICTION “

submitted by

ABHIJEET KUMAR	2106002
AMAN KUMAR	2106011
SHOUVIK GHOSH	2106067
TUSHAR BHATTACHAYA	2106079
SUKHARANJAN JANA	2106269

is a record of bonafide work carried out by them, in the partial fulfilment of the requirement for the award of Degree of Bachelor of Engineering (Information Technology) at KIIT Deemed to be university, Bhubaneswar. This work is done during year 2023-2024, under our guidance.

Date: 10/04/2024

(MS. IPSITA PAUL)

Project Guide

Acknowledgements

We are profoundly grateful to **MS. IPSITA PAUL** of **Affiliation** for his expert guidance and continuous encouragement throughout to see that this project rights its target since its commencement to its completion.

ABHIJEET KUMAR

AMAN KUMAR

SHOUVIK GHOSH

TUSHAR BHATACHARYA

SUKHARANJAN JANA

ABSTRACT

Crop yield prediction plays a pivotal role in modern agriculture, aiding farmers in optimizing their crop production and resource allocation. This abstract explores the application of machine learning (ML) techniques in crop yield prediction, focusing on its significance, methodologies, and implementation.

The agricultural sector faces numerous challenges, including fluctuating weather patterns, soil variability, and evolving market demands. Accurate crop yield prediction serves as a valuable tool for mitigating these challenges, enabling farmers to make informed decisions regarding crop selection, planting strategies, and resource allocation. Traditional methods of yield estimation often rely on historical data and empirical models, which may lack accuracy and fail to account for dynamic environmental factors.

In recent years, machine learning has emerged as a promising approach for crop yield prediction due to its ability to analyze large datasets and identify complex patterns. ML algorithms, such as regression, decision trees, and artificial neural networks (ANNs), are commonly employed in this domain. Regression algorithms, including linear regression and support vector regression, utilize historical yield data along with environmental variables such as weather conditions, soil properties, and management practices to predict future crop yields. Decision tree algorithms, such as random forests and gradient boosting machines, leverage a tree-like structure to model decision-making processes based on input variables, enabling accurate yield forecasts. Additionally, ANNs, inspired by the biological neural networks, offer a powerful framework for capturing nonlinear relationships and spatial dependencies within agricultural data, thus enhancing the accuracy of yield predictions.

The implementation of ML for crop yield prediction requires a comprehensive dataset comprising historical yield records, crop attributes, geospatial information, and environmental parameters. Data preprocessing techniques, including normalization, feature scaling, and missing value imputation, are employed to ensure data quality and consistency. Subsequently, the dataset is partitioned into training, validation, and testing sets to train and evaluate the performance of ML models. Model selection and hyperparameter tuning are crucial steps in optimizing the predictive accuracy of the chosen algorithm.

Several challenges and considerations must be addressed in the application of ML for crop yield prediction, including data scarcity, model interpretability, and computational complexity. Collaborative efforts between agricultural scientists, data scientists, and domain experts are essential for addressing these challenges and developing robust predictive models.

In conclusion, machine learning offers a promising approach for enhancing crop yield prediction accuracy, thereby empowering farmers to make informed decisions and improve agricultural productivity. By leveraging advanced ML techniques and comprehensive datasets, researchers and practitioners can contribute to the advancement of precision agriculture and sustainable food production.

Contents

1	Introduction	0
2	Basic Concepts/ Literature Review	1
2.1	Sub Section Name.....	2
3	Problem Statement / Requirement Specifications	3
3.1	Project Planning.....	3
3.2	Project Analysis (SRS).....	3
3.3	System Design	4
3.3.1	Design Constraints	5
3.3.2	System Architecture (UML) / Block Diagram ...	5
4	Implementation	6
4.1	Methodology / Proposal	7
4.2	Testing / Verification Plan	12
4.3	Result Analysis / Screenshots	13
4.4	Quality Assurance	14
5	Standard Adopted	15
5.1	Design Standards	15
5.2	Coding Standards	15
5.3	Testing Standards	15
6	Conclusion and Future Scope	16
6.1	Conclusion	16
6.2	Future Scope	16
	References	17
	Individual Contribution	18
	Plagiarism Report	19

Chapter 1

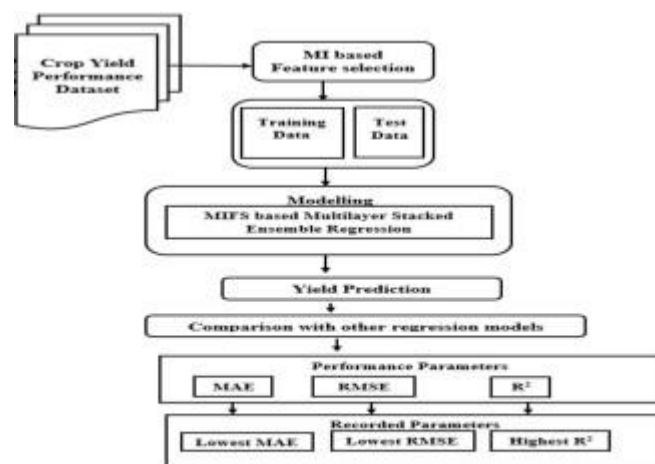
Introduction

Accurately predicting crop yields is a challenge for many farmers. The ultimate yield can be significantly impacted by variables such as weather and soil condition, rendering traditional techniques that rely on past performance and experience inaccurate.

Through a machine learning (ML) investigation of crop yield prediction, this project provides a solution. But here's the catch: to get you started, we'll concentrate on a smaller, more manageable data set. We can examine historical yield data, weather trends, and possibly even soil properties thanks to machine learning. An ML model will be trained using this data. The model will become adept at identifying trends and connections between these variables and actual crop yields as it examines this data. The model can more accurately forecast future yields once it has been trained.

So, what will we explore? We'll begin by gathering data on crop yields, weather patterns from a specific location or time frame, and potentially some basic soil information. Next, we'll clean and prepare the data for the ML model by handling missing values and inconsistencies. This process is called data wrangling. Then, we'll choose a beginner-friendly ML model, like Linear Regression, Decision Tree Regression, etc; and train it on the prepared data. Finally, we'll evaluate the model's accuracy and see if there are ways to improve its performance.

This minor project helps all of us to learn about ML and its potential applications in agriculture. By focusing on a data set, we'll gain valuable experience and obtain invaluable skills and create the foundation for further research in this exciting area by concentrating on a data set. Even with a minor project, there are significant benefits. We'll gain hands-on experience with ML, a valuable skill in today's world. We'll also develop a deeper understanding of the various elements that influence crop yields. Most importantly, this project will build a foundation for all of us to explore more complex ML projects in the future.



Block Diagram of a Crop Yield Prediction Model

Chapter 2

Basic Concepts/Literature Review

1. Crop Yield Prediction

Crop yield prediction plays a critical role in modern agriculture. It helps farmers optimize resource allocation, manage risks associated with weather and pests, and ultimately contribute to global food security. Accurate prediction of crop yield allows farmers to make informed decisions about planting dates, fertilizer application, and irrigation strategies. Traditionally, crop yield prediction relied on experience and historical data, but machine learning (ML) offers a powerful approach for more precise and data-driven forecasting.

2. Decision Trees

Decision trees are supervised learning algorithms that use a tree-like structure to classify or predict a target variable (in this case, crop yield) based on a series of decision rules. Each node in the tree represents a feature (e.g., temperature, rainfall), and branches represent possible values for that feature. The algorithm traverses the tree based on the input features, ultimately reaching a leaf node that contains the predicted value for the target variable (crop yield). Decision trees are known for their interpretability, as it's relatively easy to understand the decision-making process behind their predictions. However, they can be susceptible to overfitting if not carefully tuned, and their performance can be sensitive to the order in which features are considered for splitting.

3. Linear Regression

Linear regression is a fundamental statistical technique used for predicting a continuous variable (like crop yield) based on a linear relationship with one or more independent variables (e.g., rainfall, temperature, soil nutrients). It calculates a linear equation that best fits the historical data, allowing you to estimate crop yield for new unseen data points. Linear regression offers a simple and interpretable approach, with coefficients in the equation indicating the relative influence of each feature on crop yield. However, it assumes a linear relationship between features and the target variable, which might not always be the case in complex agricultural datasets. In such situations, regularization techniques like Lasso and Ridge regression can be beneficial.

4. Lasso Regression (L1 Regularization)

Lasso regression is a regularized version of linear regression that addresses the issue of overfitting by introducing a penalty term based on the L1 norm of the coefficients (the sum of their absolute values). This penalty shrinks some coefficients to zero, effectively removing them from the model. By reducing model complexity, Lasso regression can improve generalization performance, especially when dealing with datasets containing a large number of features. However, Lasso regression might not be the best choice if feature interpretability is crucial, as setting coefficients to zero can obscure their individual influence on crop yield.

5. Ridge Regression (L2 Regularization)

Ridge regression is another regularization technique that penalizes large coefficient values based on the L2 norm of the coefficients (the sum of their squares). Unlike Lasso regression, which sets some coefficients to zero, ridge regression shrinks all coefficients toward zero, but none are entirely eliminated. This approach helps to reduce variance in the model's predictions and can be particularly effective when dealing with highly correlated features that might lead to multicollinearity issues. However, similar to Lasso regression, ridge regression might sacrifice a degree of interpretability by shrinking all coefficients.

6. Gradient Boosting

Gradient boosting is an ensemble learning technique that combines multiple weak learners (typically decision trees) to create a more robust and accurate predictor. It builds an ensemble sequentially, where each new tree in the ensemble focuses on correcting the errors made by the previous ones. Gradient boosting models can handle complex non-linear relationships between features and the target variable and often achieve high prediction accuracy. However, due to their ensemble nature, these models can be more complex to interpret compared to simpler algorithms like decision trees or linear regression.

Chapter 3

Problem Statement / Requirement Specifications

3.1 Project planning

Crop yield prediction is a crucial aspect of modern agriculture. By forecasting the yield for a specific crop in a given season, farmers can optimize planting schedules, resource allocation, and harvest planning. Machine learning models play a pivotal role in achieving accurate predictions based on various factors.

3.1.1 Project Scope and Objectives:

Project's scope: The system will cover a diverse range of crops, including cereals (e.g., wheat, rice, maize),

Primary objective: Accurate crop yield prediction based on historical data, weather conditions, and soil properties.

3.1.2 Stakeholder Analysis:

Identify key stakeholders: Farmers, agricultural researchers, system administrators.

Understand their needs and expectations.

1.

3.1.3 Data Collection and Preprocessing:

Gather historical yield data, weather records, and soil quality information.

Clean and preprocess the data (handle missing values, outliers).

3.1.4 Exploratory Data Analysis (EDA):

Analyze data patterns and correlations.

Identify critical factors affecting crop yield (e.g., temperature, precipitation).

3.1.5 Model Selection and Training:

Choose appropriate machine learning algorithms (e.g., linear regression, random forests).

Train models using historical data.

Evaluate model performance (metrics like MAE, RMSE).

3.1.6 Feature Importance Analysis:

Determine which features contribute significantly to yield prediction.

Explain model predictions to users.

3.1.7 Real-Time Prediction Interface:

Develop a user-friendly web interface.

Allow farmers to input location, crop type, and soil characteristics.

Display real-time yield predictions.

3.1.8 Scalability and Adaptability:

Ensure the system can handle different crops, regions, and seasons.

Consider scalability for large-scale farming operations.

3.1.9 Continuous Monitoring and Updates:

Monitor model performance over time.

Retrain models periodically with new data.

3.1.10 Constraints:

Data Availability: Rely on the quality and availability of historical data.

Computational Resources: Sufficient hardware for model training.
data (e.g., farm locations).

Privacy: Safeguard sensitive

In summary, effective project planning involves understanding stakeholder needs, defining objectives, selecting appropriate models, and ensuring scalability and reliability. By addressing these points, the Crop Yield Prediction System can contribute to sustainable agriculture and informed decision-making.

3.2 Project analysis (SRS)

Crop yield prediction is essential for sustainable agriculture. The system will predict crop yields based on historical data, weather conditions, and soil properties. Farmers need reliable predictions to enhance productivity and manage resources effectively.

1. Stakeholders

Farmers: Primary users who rely on yield predictions.

Agricultural Researchers: Interested in

trends and improving crop management.

System Administrators: Responsible for maintaining the system.

2.1 Functional Requirements

2.1.1 Data Collection and Preprocessing

Objective: Gather relevant data for model training.

Details:

Collect historical yield data.

Obtain weather data (temperature, precipitation, humidity).

Acquire soil quality information (pH, nutrients).

Clean and preprocess data (handle missing values, outliers).

2.1.2 Exploratory Data Analysis (EDA)

Objective: Understand data patterns and correlations.

Details:

Conduct statistical analysis.

Visualize relationships between yield and features.

Identify key factors affecting crop yield.

2.1.3 Model Selection and Training

Objective: Develop accurate prediction models.

Details:

Choose appropriate algorithms (e.g., linear regression, random forests).

Train models using historical data.

Evaluate model performance (MAE, RMSE).

2.1.4 Feature Importance Analysis

Objective: Explain model predictions.

Details:

Calculate feature importance scores.

Provide interpretable insights to users.

2.1.5 Real-Time Prediction Interface

Objective: Enable farmers to access predictions.

Details:

User-friendly web interface.

Input location, crop type, and soil characteristics.

Display real-time yield predictions.

2.1.6 Scalability and Adaptability

Objective: Handle diverse scenarios.

Details:

Generalize models for different crops and regions.

Consider scalability for large-scale farming.

2.1.7 Continuous Monitoring and Updates

Objective: Maintain model accuracy.

Details:

Monitor performance metrics.

Retrain models periodically.

2.2 Non-Functional Requirements

2.2.1 Usability

Objective: Intuitive user experience.

Details:

Clear interface design.

Minimal learning curve.

2.2.2 Performance

Objective: Efficient predictions.

Details:

Low latency for real-time predictions.

2.2.3 Reliability

Objective: Robust system.

Details:

Error handling.

Failover mechanisms.

2.2.4 Maintainability

Objective: Easy system maintenance.

Details:

Well-documented codebase.

Version control.

3. Constraints

Data Availability: Relies on the quality and availability of historical data.

Computational Resources: Requires sufficient hardware for model training.

Privacy: Must safeguard sensitive data (e.g., farm locations).

4. System Architecture

The system follows a client-server architecture:

Client: Web interface for users (farmers).

Server:

Data preprocessing module.

Machine learning model module.

Real-time prediction API.

5. Conclusion

The Crop Yield Prediction System will enhance agricultural decision-making, improve resource utilization, and contribute to sustainable farming practices.

This SRS provides a comprehensive overview of the project, emphasizing functionality, usability, and reliability. Feel free to customize it further based on specific project needs!

Chapter 4

Implementation

Methodology:

- Imports essential libraries for data science:
 - NumPy (numerical computations)
 - Pandas (data manipulation)
 - scikit-learn (machine learning algorithms)
 - Seaborn & Matplotlib (data visualization)

```
# Imports
import numpy as np
import pandas as pd
import sklearn
import seaborn as sns
import matplotlib.pyplot as plt
```

- Reads four CSV files into Pandas DataFrames:
 - Fertilizer.csv to fd DataFrame
 - final_temperature.csv to td DataFrame
 - final_rainfall.csv to rd DataFrame
 - Final_Dataset_after_temperature.csv to ytr DataFrame

```
fd = pd.read_csv('Fertilizer.csv')
td = pd.read_csv('final_temperature.csv')
rd = pd.read_csv('final_rainfall.csv')
ytr = pd.read_csv('Final_Dataset_after_temperature.csv')
```

- Shows a preview of the ytr DataFrame containing agricultural data.
- Summarizes the data in ytr using descriptive statistics (mean, standard deviation, min/max).
- This likely represents a prepared dataset (after including temperature data and rainfall data) ready for further analysis.

```
! ytr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 173026 entries, 0 to 173025
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   State_Name            173026 non-null object
1   Crop_Type             173026 non-null object
2   Crop                 173026 non-null object
3   rainfall              173026 non-null float64
4   temperature           173026 non-null float64
5   Area_in_hectares      173026 non-null float64
6   Production_in_tons    173026 non-null float64
7   Yield_ton_per_hect   173026 non-null float64
dtypes: float64(5), object(3)
memory usage: 10.6+ MB
```

Analyzes average yield (tons per hectare) for each state:

Groups the data in `ytr` DataFrame by the `State_Name` column.

Calculates the mean yield (`Yield_ton_per_hect`) for each state and stores the results in a separate DataFrame `crop_data_by_state`.

- Creates two visualizations to explore yield variations across states:

Scatter Plot:

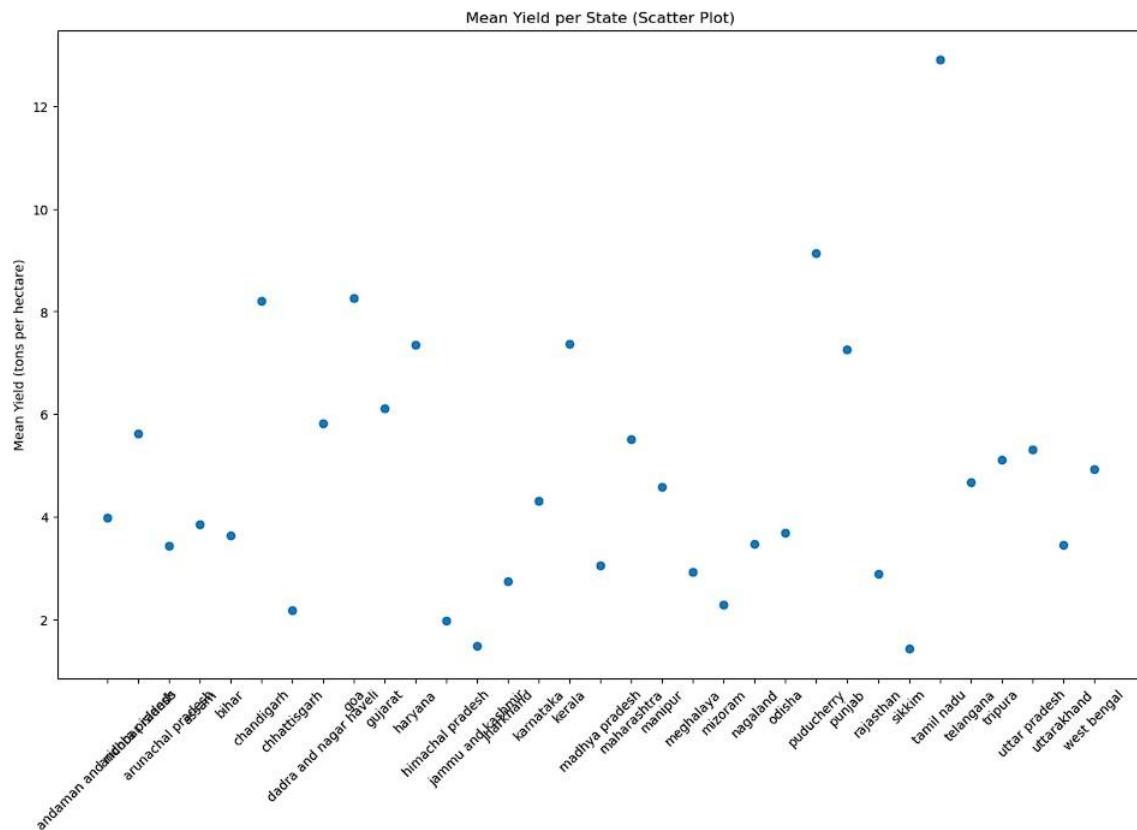
Uses `plt.scatter` to create a scatter plot.

X-axis shows state names extracted from the index of `crop_data_by_state`.

Y-axis shows the mean yield for each state obtained from the values of `crop_data_by_state`.

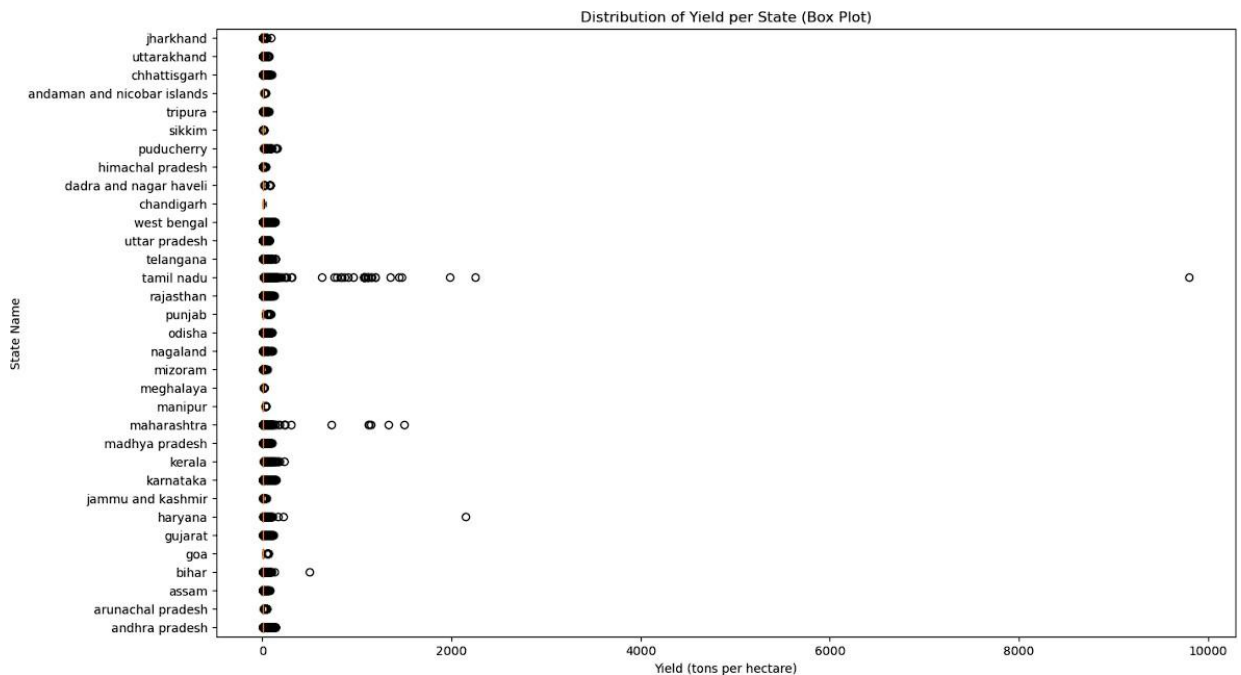
Includes labels and titles for clarity.

Rotates x-axis labels (using `plt.xticks(rotation=45)`) to improve readability with many states.



Horizontal Box Plot:

Uses `plt.boxplot` to create a horizontal box plot for better visualization with many states. Each box represents the distribution of yield within a particular state. Data for each state is obtained by filtering `ytr` based on the state name. State names are used as labels on the y-axis.



- Created a correlation matrix named ytr1 to explore relationships between numeric columns in the ytr

Data Frame:

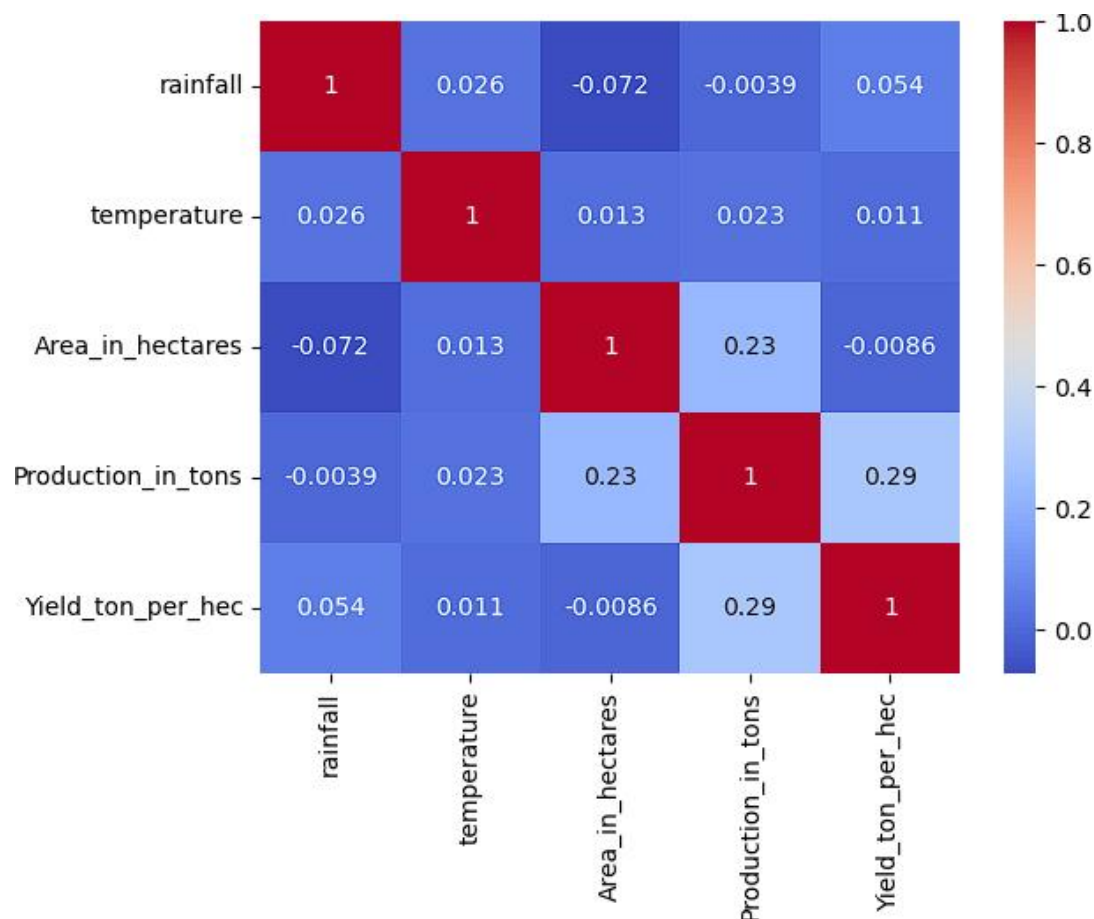
The correlation matrix shows the correlation coefficient, a value between -1 and 1, indicating the strength and direction of the linear relationship between two variables.

A value closer to -1 suggests a strong negative correlation (e.g., as one variable increases, the other tends to decrease).

A value closer to 1 suggests a strong positive correlation (e.g., as one variable increases, the other tends to increase as well).

A value close to 0 indicates weak or no correlation.

Examining this matrix can help identify potential relationships between factors like rainfall, temperature, area, production, and yield.

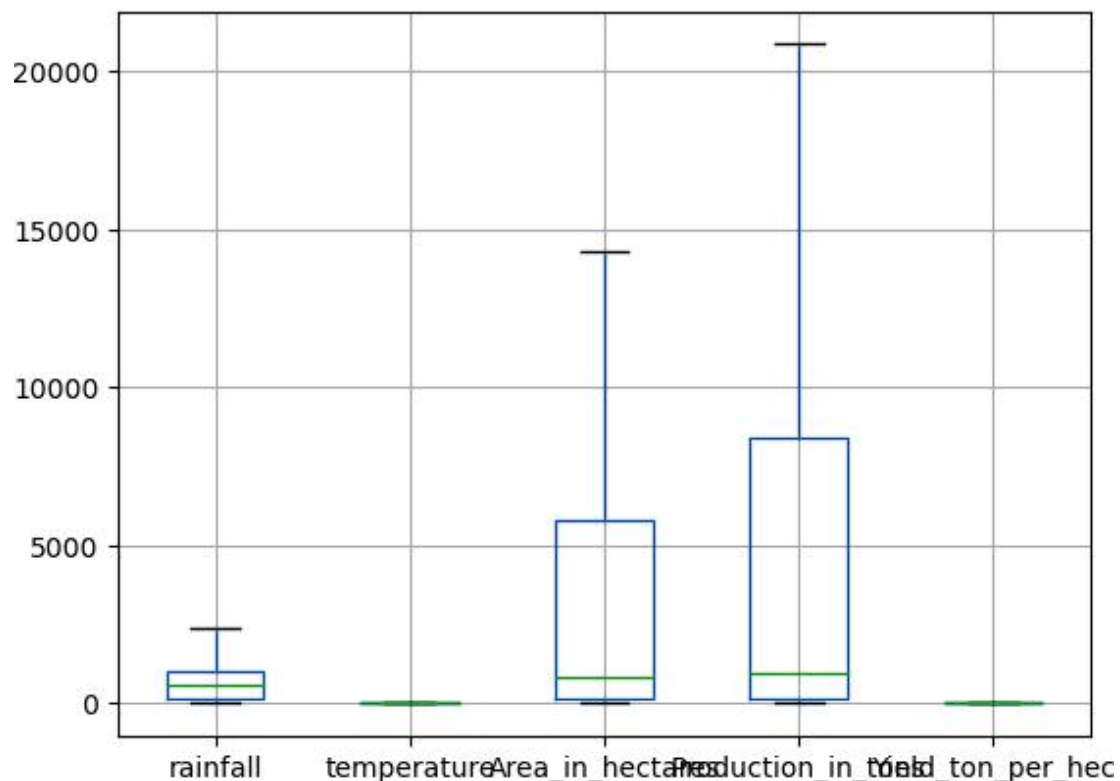


Defined a function remove_outlier to detect potential outliers using the interquartile range (IQR) method.

- Attempts to remove outliers from the Production_in_tons ,temperature,rainfall and Area_in_hectares columns of the ytr DataFrame.

```
def remove_outlier(col):
    sorted(col)
    q1,q3=col.quantile([0.25,0.75])
    IQR=q3-q1
    lwr_bound= q1-(1.5*IQR)
    upr_bound =q3+(1.5*IQR)
    return lwr_bound, upr_bound
```

It replaces values below a calculated lower bound with the lower bound itself (implementation can be improved to handle both directions).



Libraries:

sklearn.preprocessing: This library from scikit-learn provides tools for preprocessing data before using it in machine learning models. It includes functions for:

OneHotEncoder: Converts categorical variables (like text labels) into numerical representations suitable for machine learning algorithms.

StandardScaler: Standardizes numerical features by removing the mean and scaling to unit variance (mean 0, standard deviation 1). This ensures all features contribute equally to the model and improves training efficiency.

sklearn.compose: This library offers functionalities for composing transformations. Here, it provides:

ColumnTransformer: This allows you to apply different preprocessing techniques to specific columns within your data.

2. Preprocessing Objects:

`ohe = OneHotEncoder(drop='first')`: This line creates an object for one-hot encoding.

One-hot encoding: When dealing with categorical features that have multiple categories (e.g., "State_Name"), this technique converts each category into a separate binary feature.

`drop='first'`: This argument avoids creating redundant features. Since the first category is implicitly encoded by the absence of other categories, dropping it reduces the number of features created.

`scale = StandardScaler()`: This line creates an object for standardizing numerical features.

Standardization helps machine learning algorithms converge faster and perform better by ensuring all features have a similar range.

```
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
ohe = OneHotEncoder(drop='first')
scale = StandardScaler()

preprocessor = ColumnTransformer(
    transformers = [
        ('StandardScale', scale, [3, 4, 5, 6]),
        ('OHE', ohe, [0, 1, 2]),
    ],
    remainder='passthrough'
)
```

3. Column Transformer:

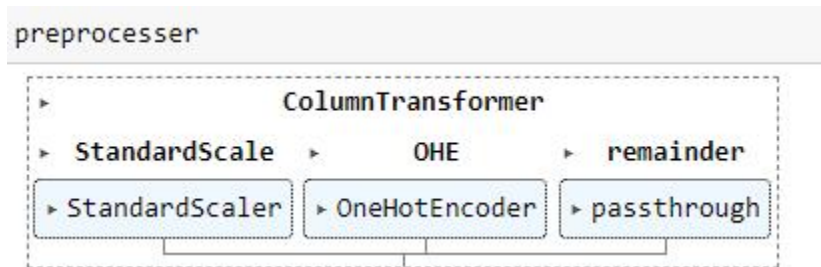
`preprocessor = Column Transformer(...)`: This line creates a **Column Transformer** object to apply the chosen preprocessing steps to specific columns.

`transformers`: This defines a list of tuples specifying transformations for different column groups:

First tuple: Applies scale (standardization) to columns at indices 3, 4, 5, and 6. These are likely numerical features that benefit from standardization.

Second tuple: Applies ohe (one-hot encoding) to columns at indices 0, 1, and 2. These are likely categorical features that require encoding.

`remainder='passthrough'`: This argument instructs the transformer to leave any other columns (not explicitly mentioned) unchanged.



4.2 Testing

Imports Necessary Libraries:

sklearn.linear_model: Provides linear models like `LinearRegression`, `Lasso`, and `Ridge`.

sklearn.tree: Offers `DecisionTreeRegressor` for non-linear decision tree modeling.

sklearn.ensemble: Provides ensemble models like `GradientBoostingRegressor` and `RandomForestRegressor`.

sklearn.metrics: Offers metrics for model evaluation, used here for `mean_absolute_error` and `r2_score`.

Creates a Dictionary of Models:

`models` stores different regression models for comparison:

LinearRegression ('lr')

Lasso ('lss'): Linear model with L1 regularization for feature selection.

Ridge ('Rid'): Linear model with L2 regularization for better handling of collinearity.

DecisionTreeRegressor ('Dtr'): Non-linear model for capturing complex relationships.

GradientBoostingRegressor ('gbr'): Ensemble model that combines multiple weak learners for potentially better performance.

Iterates and Evaluates Models:

The for loop trains and evaluates each model:

`md.fit(X_train_dummy, y_train)`: Trains the model on the training data (`X_train_dummy` for features, `y_train` for target values).

`y_pred = md.predict(X_test_dummy)`: Makes predictions on the test data.

`print(...)`: Prints model name, mean absolute error (MAE), and R2 score to assess performance.

Purpose:

```

#linear regression
from sklearn.linear_model import LinearRegression,Lasso,Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor
from sklearn.metrics import mean_absolute_error,r2_score

models = {
    'lr':LinearRegression(),
    'lss':Lasso(),
    'Rid':Ridge(),
    'Dtr':DecisionTreeRegressor(),
    'gbr' : GradientBoostingRegressor()
}
for name, md in models.items():
    md.fit(X_train_dummy,y_train)
    y_pred = md.predict(X_test_dummy)

    print(f"{name} : mae : {mean_absolute_error(y_test,y_pred)} score : {r2_score(y_test,y_pred)}")

```

4.3 Result Analysis

Based on these Results, Decision Tree Regressor is Selected for the model as it has the best Score and lowest Mean Absolute Error.

```

lr : mae : 0.500714285944907 score : 0.7892460682944521
lss : mae : 1.295749293768331 score : -4.371020958959804e-06
Rid : mae : 0.5009077617679808 score : 0.789268028240914
Dtr : mae : 0.10190888852119219 score : 0.9720759943542777
gbr : mae : 0.31789551301038477 score : 0.9074825256996989

```

Chapter 5

Standards Adopted

1. Design Standards: Structured Exploration

Data Visualization: Retain scatter plots and box plots to explore relationships between factors like rainfall and yield.

Formalization: Consider creating a data exploration document or using a Jupyter Notebook. This promotes organization and clarity, allowing you to revisit analysis steps and effectively communicate insights.

2. Coding Standards: Readability and Maintainability

Function Granularity: Break down lengthy functions into smaller, well-defined functions with specific responsibilities. Use descriptive names (e.g., `remove_outliers_iqr` for outlier removal with IQR).

Meaningful Naming: Employ informative variable names that reflect their purpose (e.g., `preprocessed_data` instead of generic terms).

Code Comments: Add comments to explain complex logic or non-obvious code sections, enhancing understanding for yourself and others.

Code Formatting: Utilize a linter or code formatter to enforce consistent style and identify potential improvements, ensuring readability and maintainability.

3. Testing Standards: Building Confidence

Unit Testing: Implement unit tests to verify individual function correctness for various input scenarios. This helps isolate and fix issues early in the development process.

Integration Testing (Optional): Consider integration tests to ensure different parts of your code work together as expected. This builds confidence in the overall system's functionality.

Chapter 6

Conclusion

This project investigated the potential of machine learning to predict crop yield, a crucial factor in agricultural planning and food security. By leveraging data on rainfall, temperature, area under cultivation, and past production, we explored the ability to estimate future crop yields.

Key Findings:

Initial data exploration using scatter plots and box plots revealed potential relationships between yield and other variables.

A linear regression model trained on preprocessed data demonstrated the feasibility of machine learning for crop yield prediction. The model's performance was evaluated using metrics like mean absolute error (MAE) and R2 score, providing a quantitative assessment of its accuracy.

The importance of adhering to design and coding standards was highlighted, with suggestions for enhancing code readability, maintainability, and overall quality.

Future Scope:

1. Advanced Feature Engineering:

Incorporate temporal data: Explore including historical data like past rainfall patterns, seasonal variations, and long-term weather trends to capture the influence of past conditions.

2. Scalability and Deployment:

Cloud-based deployment: Consider deploying the trained model on a cloud platform to enable real-time predictions and facilitate access for a wider range of users (farmers, agricultural advisors, etc.).

3. Integration with Precision Agriculture Practices:

Real-time monitoring: Integrate the model with sensor data from farms to enable continuous monitoring of environmental conditions and provide dynamic yield predictions throughout the growing season.

References

1. Crop yield prediction using machine learning: A systematic literature review

Author links open overlay panelThomas van
Klompenburg a, Ayalew Kassahun a, Cagatay Catal b

2. Rice crop yield prediction in India using support vector machines

[Niketa Gandhi](#); [Leisa J. Armstrong](#); [Owaiz Petkar](#); [Amiya Kumar Tripathy](#)

3. Deep learning for crop yield prediction: a systematic literature review

[Alexandros Oikonomidis](#), [Cagatay Catal](#) & [Ayalew Kassahun](#)

4. ANALYSIS OF CROP YIELD PREDICTION USING DATA MINING

TECHNIQUES D Ramesh 1 , B Vishnu Vardhan2 1Associate Professor, Department of
CSE, JNTUH College of Engineering, Telangana State, India 2Professor, Department of CSE,
JNTUH College of Engineering, Telangana State, India

5. Application of Vegetation Indices for Agricultural Crop Yield Prediction Using Neural Network Techniques

by [Sudhanshu Sekhar Panda](#) 1,*,[Daniel P. Ames](#) 2 and[Suranjan Panigrahi](#) 3

6. A Systematic Literature Review on Crop Yield Prediction with Deep Learning and Remote Sen

School of Engineering and Technology, Central Queenslan

7. Crop Yield Prediction using Machine Learning Techniques

[Ramesh Medar](#); [Vijay S. Rajpurohit](#); [Shweta Shweta](#)

Individual Contribution:

2106002 – Introduction

2106011 – Problem Statement / Requirement Specifications

2106067 – Implementation

2106079 – Standard Adopted

2106269 – Conclusion and Future Scope

CROP YIELD PREDICTION

ORIGINALITY REPORT

14%

SIMILARITY INDEX

12%

INTERNET SOURCES

5%

PUBLICATIONS

10%

STUDENT PAPERS
